

Machine learning text classification

Colab link:

https://colab.research.google.com/drive/1O6u1s2JaaVOZ5MiFot_v6Zn7F2ujv_5n

1) Data preprocessing:

Because the data was encrypted, I decided not to change characters format. For example “A@” will be treated as different from “a@”

No non-literary character is eliminated.

2) Features extraction:

To get numeric values from the text I used tf-idf, the default version from sklearn: TfidfVectorizer.

The n-grams gave a better accuracy than the simple words, the most efficient variant was the 6-character n-grams. I used the ‘char_wb’ parameter to only get the n-grams inside words.

To normalize the data I used l2.

N-grams

Full words

	precision	recall	f1-score	support
0	0.69	0.75	0.72	1301
1	0.74	0.68	0.71	1355
accuracy			0.71	2656
macro avg	0.72	0.72	0.71	2656
weighted avg	0.72	0.71	0.71	2656

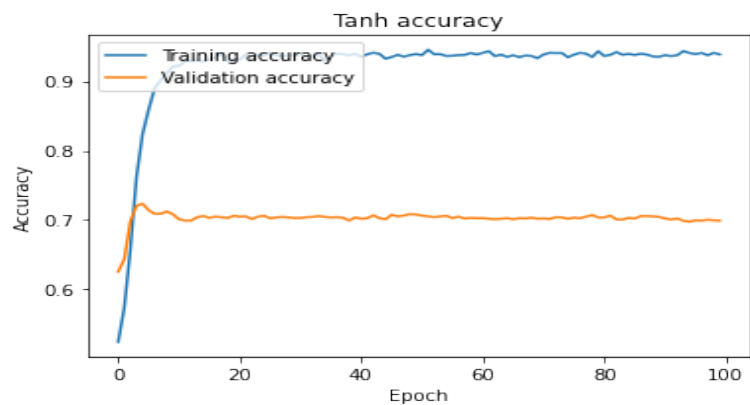
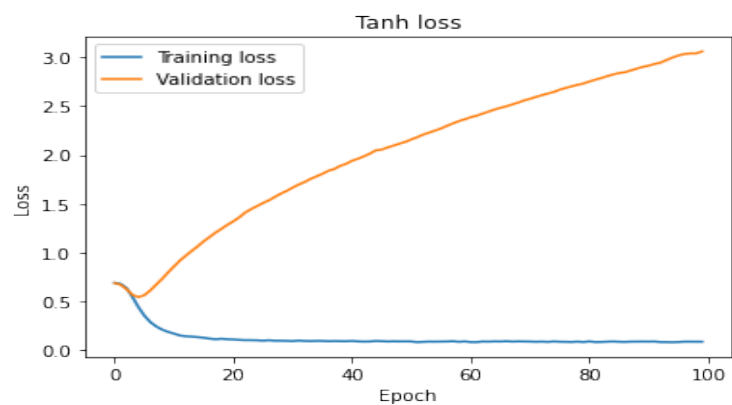



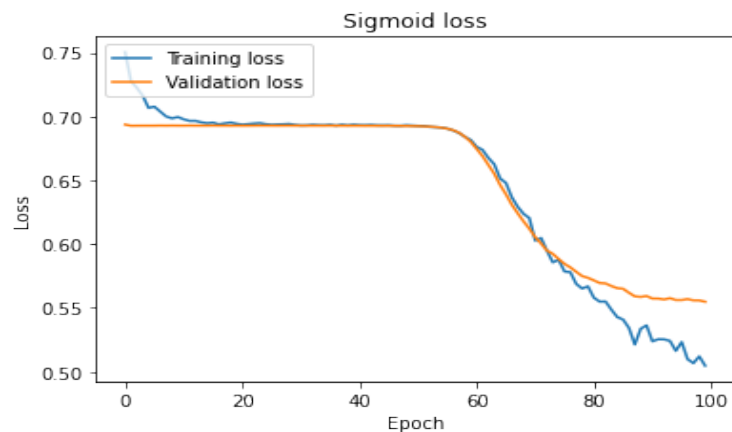
	precision	recall	f1-score	support
0	0.64	0.71	0.67	1301
1	0.69	0.61	0.65	1355
accuracy			0.66	2656
macro avg	0.66	0.66	0.66	2656
weighted avg	0.67	0.66	0.66	2656

3) Neural network model:

I used keras to easily implement a mini network with two hidden layers. With a small number of perceptrons per layer I had good results,

otherwise the model would be overfitting very quickly. The optimal number of epochs I have found is 100.

For the last layer I used the sigmoid activation function, for the rest I tried several variants.



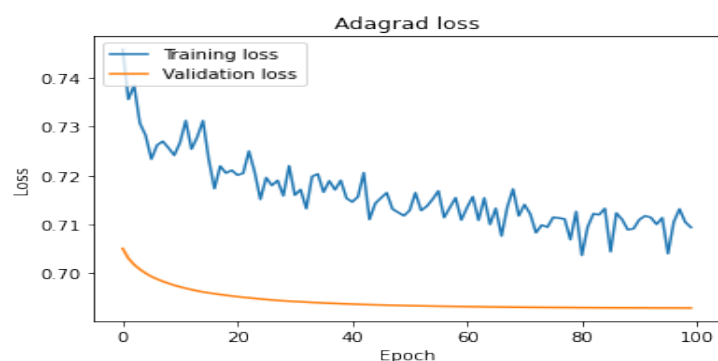

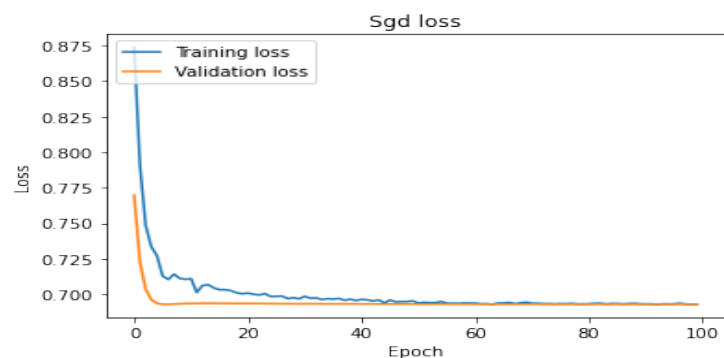


As can be seen from the graphs, the best model was the one with sigmoid.

I've also used a dropout layer (with a value of 0.5) to prevent overfitting.

As for the loss function most of those offered by keras had similar performance, but the best was binary_crossentropy.

The situation is different for gradient descent functions. Sgd and adagrad both are having a poor performance both on training and on validation sets, on the other hand adam gives good results.



Model Summary

Model: "sequential"		
Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 4)	251396
activation (Activation)	(None, 4)	0
dropout (Dropout)	(None, 4)	0
dense_1 (Dense)	(None, 8)	40
activation_1 (Activation)	(None, 8)	0
dropout_1 (Dropout)	(None, 8)	0
dense_2 (Dense)	(None, 4)	36
activation_2 (Activation)	(None, 4)	0
dropout_2 (Dropout)	(None, 4)	0
dense_3 (Dense)	(None, 1)	5
activation_3 (Activation)	(None, 1)	0
=====		
Total params: 251,477		
Trainable params: 251,477		
Non-trainable params: 0		
=====		

Classification Report

	precision	recall	f1-score	support
0	0.76	0.61	0.68	1301
1	0.69	0.82	0.75	1355
accuracy			0.72	2656
macro avg	0.73	0.72	0.71	2656
weighted avg	0.73	0.72	0.72	2656

Confusion Matrix

[[800 501]	
[247 1108]]	